


1. Collect and plot running times in order to answer each of the following questions:


- Is the running time of the `addFirst` method $O(c)$ as expected? How does the running time of `addFirst` for `DoublyLinkedList` compare to `add(0, item)` for `ArrayList`?



- Yes, the running time for `addFirst` was indeed $O(c)$. As you can see in the `ArrayList`'s `add()` timings, while it is $O(c)$ as well, it is much larger since it has to move over all of the currently present elements first.
- Is the running time of the `get` method $O(N)$ as expected? How does the running time of `get(i)` for `DoublyLinkedList` compare to `get(i)` for `ArrayList`?



- Yes, `get()` was $O(N)$ as shown in the graph, while `ArrayList`'s was $O(c)$ due to its ability to directly look up values based on index.
- Is the running time of the `remove` method $O(N)$ as expected? How does the runtime of `remove(i)` for `DoublyLinkedList` compare to `remove(i)` for `ArrayList`?



- Yes, `remove` was $O(N)$, as was `ArrayList`'s. The `ArrayList` version was much faster, though, since it can go straight to the correct index, even though it had to copy everything back.
- In general, how does `DoublyLinkedList` compare to `ArrayList`, both in functionality and performance?
 - A `DoublyLinkedList` isn't directly indexed, so it has to cycle through the entire list to get to the point it needs to look at, making `get` take much longer than an `ArrayList`. However, since each element of the list is separate with simple pointers showing what's "next", and overwriting those is a simple call, adding into or removing out of the list is much faster, as `ArrayList`'s indexing must be almost completely adjusted every time a new element is added or removed.
 - In general, how does `DoublyLinkedList` compare to Java's `LinkedList`, both in functionality and performance?
 - A traditional `LinkedList` only has a pointer to the beginning element of the list, followed by a long, one way chain, making reaching the final element potentially very expensive. However, with the `DoublyLinkedList`, we also have access to the end of the List, not only making reaching the end an instantaneous constant time function, but now reduces the furthest we have to search in an N element array from all N elements down to $N/2$.
 - Compare and contrast using a `LinkedList` vs an `ArrayList` as the backing datastructure for the `BinarySearchSet` from Assignment 3. Would the Big-O complexity change on `add`, `remove`, and `contains`?
 - Yes, and for the worse for all three, as the reason the `BinarySearchSet` worked so well was because we could easily skip around, from midpoint to midpoint, cutting down the search size by half until we narrowed it down to where the element should be. This took huge advantage of the indexes within the `ArrayList`. However, with our current

DoublyLinkedList, even if we had a constantly adjusting pointer to the center of the overall list, we would still have to step through every element just to get to the middle to make a new comparison. Since we would probably pass the element we were searching for on the way anyways, it'd be better to simply search through the list and check each element as we go, which is what we do at the moment anyways with our current get function.

5. How many hours did you spend on this assignment?
 - I spent about 5 - 6 hours on this assignment total.